Word Embeddings

Intro to neural language modeling

Outline

1. Key ideas in natural language
2. Classic natural language ideas
3. Word embeddings
4. Neural word embeddings
   a. GloVe
   b. Word2Vec + Skipgram
5. Practicalities
6. Additional resources

What is one-hot encoding (OHE)?

OHE is the basis of classic ML in NLP

It is a method that converts categorical data into a numeric vector that computers can understand and is easy to interpret

The way it works is by constructing a V length vector which is the length of the vocabulary (and where each index represents a word)

For a document that has a word a 1 is assigned to that index in the vector

One hot vectors can be very large depending on the vocabulary

Frequently it is used with the "hashing trick" and "sparse encoding" to reduce the vector size

DOCUMENT SPACE

DOCUMENT 1: THE QUICK BROWN FOX
DOCUMENT 2: THE LAZY BROWN DOG

VOCABULARY

{ THE, QUICK, BROWN, FOX, LAZY, DOG }

W_{quick} = [0, 1, 0, 0, 0, 0]

ONE-HOT ENCODED REPRESENTATION: QUICK

A simple language models uses matrices to represent relations among words/documents

- Language can be represented through:
    - Term-document (TDMs) or
    - Document-term matrices (DTMs - transpose of TDMs)
- A document is the sum of one-hot-encoding word-vectors
- We frequently call this the *bag-of-words representation*

- The relation among documents is based on some metric that assesses the difference/similarity between document vectors

- Models that use DTMs are called <u>vectorial-semantic</u> models (VSMs)

DOCUMENT 1

$$W_{the} + W_{quick} + W_{brown} + W_{fox}$$
$$=$$
$$[\ 1, 1, 1, 1, 0, 0\ ]$$

DOCUMENT ENCODING

Summing OHE elements gives us a BOW vector

	DOCUMENT 1	DOCUMENT 2
THE	1	1
QUICK	1	0
BROWN	1	1
FOX	1	0
LAZY	0	1
DOG	0	1

TERM DOCUMENT MATRIX

*Encoding two different documents - notice that sequence is obfuscated in BOW*

Classic ML NLP approaches remain important

Machine learning methods that use DTMs include:
- Search
- Topic/document clustering
- Word disambugation
- Sentiment analysis
- Automatic key phrase extraction

Note:
- Some recent developments use a mixture of classic models and neural embeddings such as LDA2Vec
- Many deep NLP methods build on foundational concepts in NLP
- Studying classic approaches and their problems gives you insight into how neural approaches resolves them
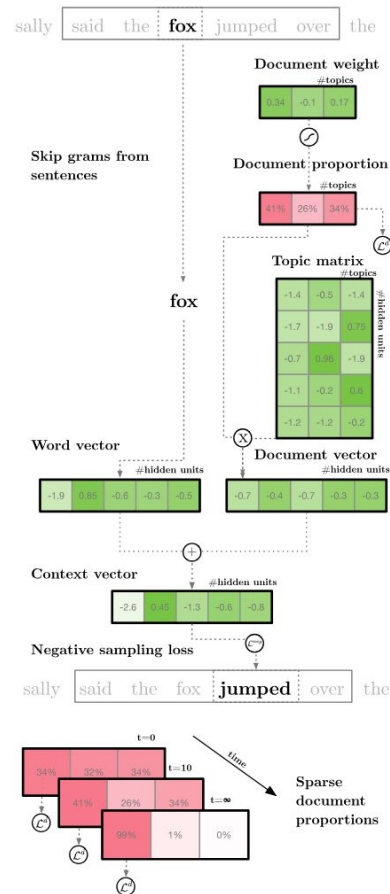
** Source: Mixing Dirichlet Topic Models and Word Embeddings to Make lda2vec
https://www.groundai.com/project/mixing-dirichlet-topic-models-and-word-embeddings-to-make-lda2vec/

What are some OHE of limitations?

OHE does not consider relations among words

OHE is used to measure statistical properties across phrases and corpora in extremely high dimensions

By definition OHE assumes words are orthogonal to one another
- For example the words: _pencil_ and _pen_ are independent of one another although they are both are writing instruments
- Relations among words only surface as a result of the statistical similarities among global structure in documents
- This is a central limitation since local information about sequence and position inform the relation among words

Pencil and pen are treated as orthogonal despite being writing instruments when using OHE

Topics in space

Term Space

iphone

iphone Fans

fans

baseball

Sports Fans w/ Cameras

<number>
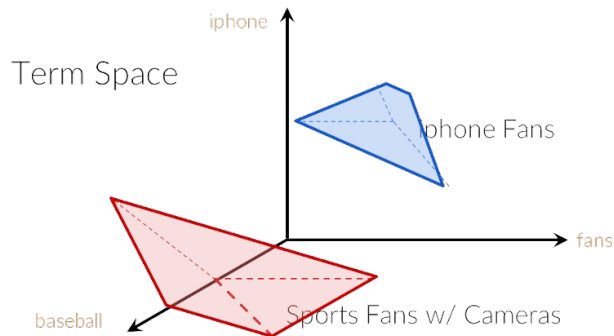
Vector-semantic models: Documents are reconstructed in term space through the addition of OHE, notice that baseball, fans and iphone are orthogonal to one another in 3-space

Reducing Dimensionality of the System

Early attempts to reduce the dimensionality of systems and generate lower dimensional representations of the DTM included SVD

SVD is a generalization of PCA to non-square matrices

It is a decomposition technique that uses "singular" rather than "principal" values to approximate the DTM

Early applications included latent semantic analysis/indexing for recommendations systems (Netflix prize) and is still used in search (Google banners)

While the original reconstruction error was built to optimize user-movie relations, re-focusing on local term-term relations we are able to generate meaningful embeddings and is the basis of GloVe(more to come on this later)

SVD through a user-movie matrix

Singular-Value decomposition:

$$X_{m \times n} = U\Sigma V^T$$

where:

- $U_{m \times d}$ - latent user matrix
- $V^T_{d \times n}$ - latent movie matrix
- $\Sigma_{d \times d}$ - Singular value matrix.

Example: User- Ratings:

Ratings matrix: $X = U_{6 \times 5} \times \Sigma_{5 \times 5} \times V^T_{5 \times 5}$

Reconstruction error using MSE:

$$\text{argmin}||X - X_k||_F = \text{argmin } \Sigma_i \Sigma_j (X(i,j) - X_k(i,j))^2$$

$$= \underset{U,V,\Sigma}{\text{argmin}} \Sigma_i \Sigma_j (X(i,j) - U(i,.) \times \Sigma_k \times V(.,j))^2$$

		Alien	Predator	Matrix	Casablanca	Amelie
SCI-FI	USER 1	4	4	3	0	0
	USER 2	5	4	4	0	0
	USER 3	2	4	5	0	0
	USER 4	5	5	3	0	0
ROMANCE	USER 5	0	1	0	5	4
	USER 6	0	1	0	4	5

USER-MOVIE MATRIX (X)

Example: User-Ratings

Clearly there are two genres of movies, along with differential ratings for users.

How could we decompose this matrix to provide recommendations and approximate its reconstruction like in PCA?

We can't use PCA because the matrix isn't square so we use SVD

SVD in Pictures - Lower dimension user space
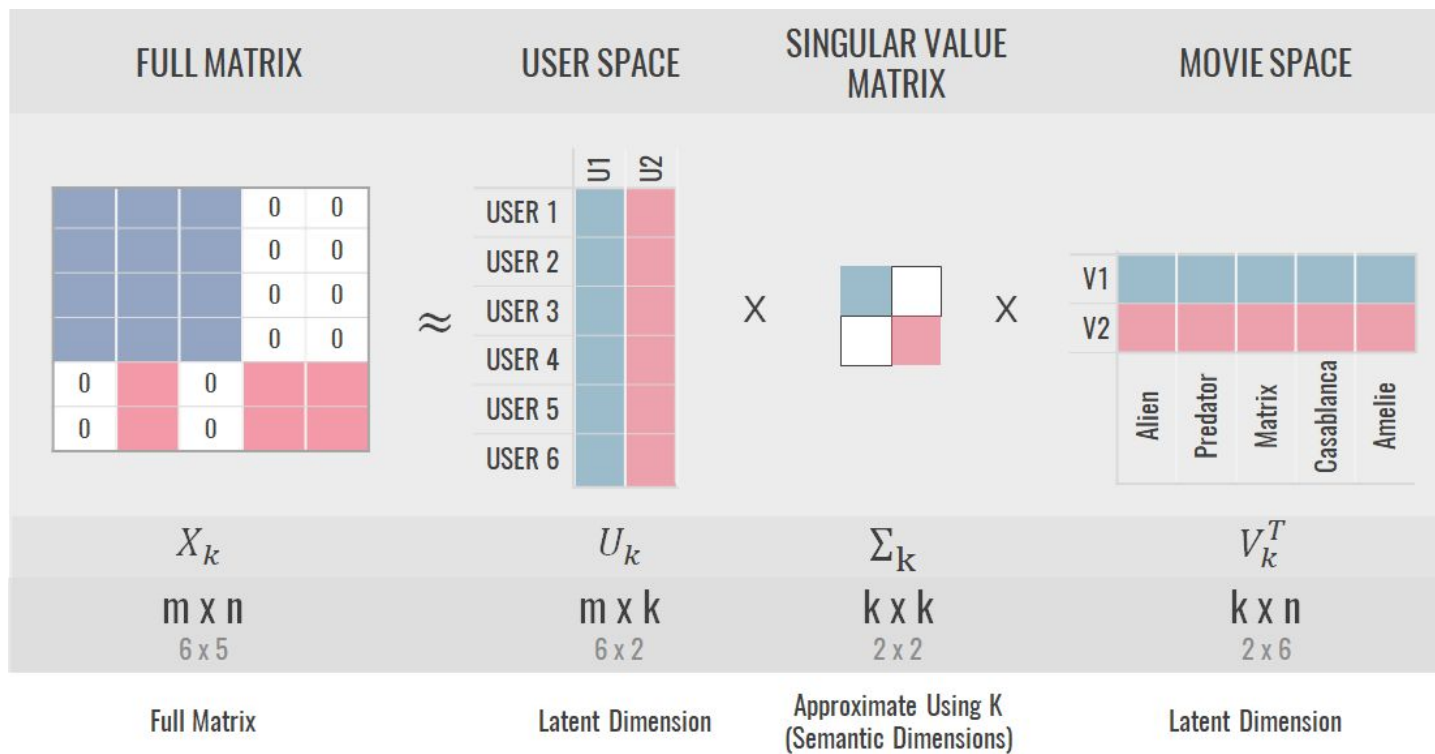
The rows in the user space X approximates X whole rank 2 in V^T_2 exist in R

Measuring similarity in semantic space

Word i = u_i

Document i = v_i

u_i u_j

query vectors

Latent (Term) Space

D = 5, k = 2

SIMILARITY CALCULATION

Term:
$$\text{sim}(w_i, w_j) = \cos(u_i \Sigma_k, u_j \Sigma_k)$$

Document:
$$\text{sim}(d_1, d_2) = \cos(\Sigma_k v_i, \Sigma_k v_j)$$

0 ← — — — — — — $\cos(x_1, x_2)$ — — — — — — → 1

Not Similar   Similar

$$\cos(x_1, x_2) = \frac{x_1 \cdot x_2}{||x_1|| \times ||x_2||}$$

Recall from linear algebra that this can be viewed as an orthogonal projection of one vector onto another.

Note the cos similarity is a value between -1 and 1.

A value of cos(x1,x2) = 0 indicates the vectors are orthogonal and of -1 that they are in opposite directions

The normalized direction gives us an indication of how close the vectors are through it's angle

An alternate view

Since singular values are distributed in the similarity query between two users or movies, recommendation systems are frequently constructed using only U (user) and V (movie)

The reconstruction error is then written as:

$\text{argmin}||X - X_k||_F = \underset{U,V}{\text{argmin}} \ \Sigma_i \Sigma_j (X(i,j) - U(i,.)V(.,j))\text{^2}$

The interpretation views the user and movie ratings as a linear combination

This looks very similar to the minimization problem for NNs

We can use **SGD** to minimize for the values of U and V or **alternating least squares**

Linear model with 2 biases   Penalties on u and v

$$L = \Sigma_i \Sigma_i (X(i,j)$$
$$-\boldsymbol{u}_i^T \cdot \boldsymbol{v}_j + b_i + c_j)^2 + \lambda_u \Sigma_i \ // \ \boldsymbol{u_i} \ // \ ^2 + \lambda_v \Sigma \ // \ \boldsymbol{v}_j \ // \ ^2$$

We can also add in bias terms and a regularizers and minimize with respect to U,V

For details on factorization see: https://www.ethanrosenthal.com/2016/01/09/explicit-matrix-factorization-sgd-als/

Word embeddings in the real world

Word embeddings are used in several common applications, including:

1. Feature extraction
2. Sentiment analysis
3. Question answer
4. Translation
5. Anywhere you would use language in your pipeline
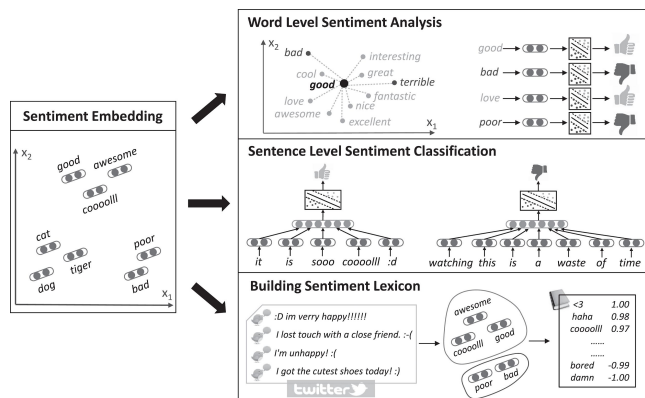
Why use word embeddings?

Pros of using Embeddings	Caveats when using Embeddings

Word embeddings bridge the "meaning gap"
- Words in relation to one another have meaning and are no longer orthogonal

Reduces the size of the object contained in memory
- Generates a low dimensional (between 50-300) embedding of words
- Using a DTM with 50K words and 10^6 documents would be roughly 190Gigabytes (stored as float32s)

More expressive representation of the data
- Analogical reasoning enables vector math with words
- Composition of word vectors is more interpretable

- Positions in space of word-vectors do not hold significance in that space/dimension
    - Only their relative distances/angles to one another do
- The embeddings are not unique
- There are different embedding approaches two common ones are
    - Probabilistic prediction approaches (predict based on context/center words) and
    - Reconstruction approaches (approximate the DTM)
- Generating good embeddings requires a lot of data and training them is challenging

Embeddings

Instead of approximating the TDM we embed words as a set of vectors in some high dimensional space (although still lower than TDM space)

There is more than enough information in these high dimensions to express the relation among words (usually D=50,100,300 - empirically after 300 dims there is no gain)

On the right, the example, illustrates several type relations, note the angle, direction and trajectory

A common example is:
vec(king)−vec(man)+vec(woman)≈vec(queen)

However other like grammar, capital among countries and compositions (like the one above) similarly exist
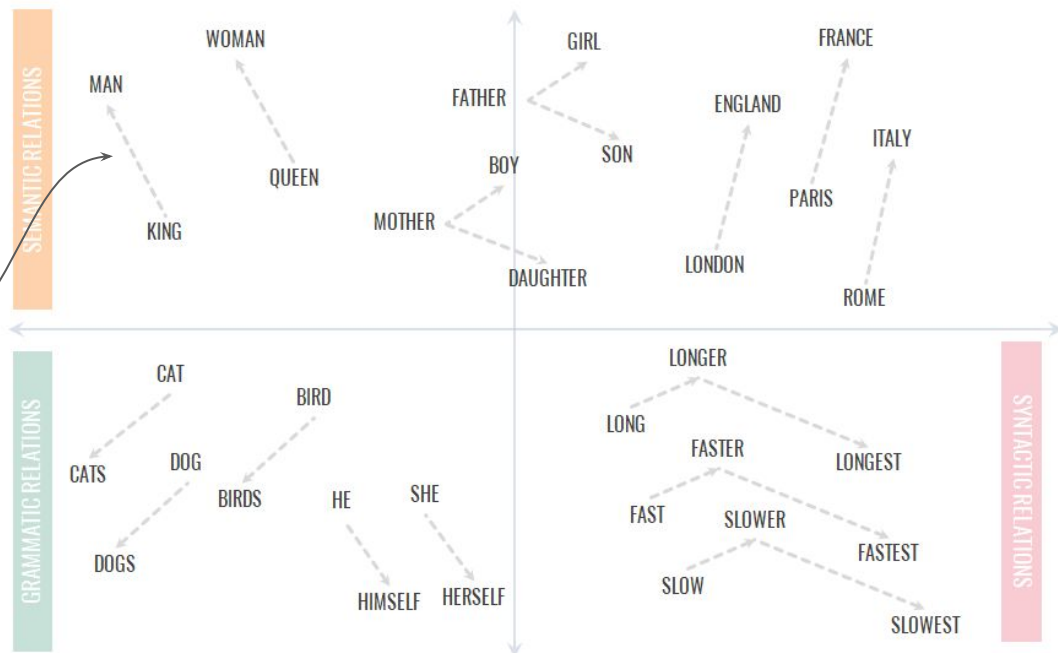
Figure: Example 2D Embedding Space Obtained Through Word2Vec

Embeddings enable vector math

The relations among embeddings can be summarized using vector math - summing and subtracting allows us to interpret relations, i.e.:

vec(king)−vec(man)+vec(woman)≈vec(queen)

The distances and angles between vectors therefore can be interpreted as "similarity of idea/structure"

Using this approach we use analogic reasoning to approximate similar relations between pairs of words


2d Projected Vector Embeddings

Semantic Relationships

MAN:WOMAN :: KING: ?

MAN:WOMAN :: KING: ?

2d Projected Vector Embeddings

$$x_{woman} - x_{man} + x_{king}$$
$$=$$
$$x_{king} - x_{man} + x_{woman}$$

$$d = \text{argmax}_i \frac{(x_{woman} - x_{man} + x_{king})^T x_i}{||x_{woman} - x_{man} + x_{king}||}$$

Common Embedding approaches

Two common approaches to embedding are:
1. GloVe - Global Vector (reconstructive) and
2. Word2Vec (predictive), which is made of two models
   a. Continuous bag of words and
   b. Skipgram

We'll focus first on GloVe and then discuss Word2Vec

GloVe and Skipgram type models tend to be most used in real world applications
- GloVe - similar idea to rec systems
- Skipgram - easy application

CBOW predicts a center word from context words

Skipgram predicts context words from a center word

GloVe

A reconstructive approach

Uses a weighted term-term matrix (call it X)

X is constructed such that i is the center word and j is the context word

Context distance is calculated as distance from one another

Example: "The quick brown fox jumped"
- X(The, quick) = +1
- X(The, brown) = +½
- X(The, fox) = +⅓

Implications include that X is a sparse matrix and will have a similar to solution to rec system

We minimize the loss below with respect to the weights in vector $U, V, b\ and\ c$. The resulting approach generates a set of embeddings U, V

f is a weighting function

1 is added in because X is sparse

$$L=\sum_i\sum_j f(X(i,j))(\log(X(i,j)+1)$$
$$-u_i^T \cdot v_j + b_i + c_j)^2 + \lambda_u \parallel U \parallel^2 + \lambda_v \parallel V \parallel^2$$

We reconstruct the log of the term-term matrix - we can do this because log is a bijective mapping

U and V are two alternate word embeddings

Setup: Skipgram

Skipgram is part of the Word2Vec probabilistic predictive embedding models

In Word2Vec, local relations provide greater contextual information than global structure

Accordingly two important ideas for Word2Vec models are "center" and "outer context" words

Context is preserved by calculating blocks of probability for surrounding words with the skip gram models

The probability distribution looks like

$$Pr(w_{t-m},...,w_{t-1},w_{t+1},...,w_{t+m}|w_t;\theta) = \prod_{-m \leq j \leq m; j \neq 0} Pr(w_{t+j}|w_t;\theta)$$

The figure and table illustrate center-context word dependence for a window of m=3

Example:
"The quick brown fox jump over the lazy dog"

For a context window of m = 2 and target of jumps we have the following target-context probability for one phrase

$$\prod_{-m \leq j \leq m; j \neq 0} Pr(w_{t+j}|jumps;\theta)$$
$$=$$
$$Pr(u_{brown}|v_{jumps})Pr(u_{fox}|v_{jumps})Pr(u_{over}|v_{jumps})Pr(u_{the}|v_{jumps})$$

Architecture: skipgram

The skipgram architecture uses two matrices V and U to calculate the embeddings of the words

The matrices are optimized against a softmax function with cross entropy error

The input is a word and the output is the set of context words probability

The architecture can be trained using softmax however is very slow, several alternate methods have been proposed including "negative sampling"

Summary:
- Input: Target word
- Output: Context words
- Output layer: softmax
- Loss: Cross entropy

$$Pr(w_{quick}|w_{fox};\theta)$$

Skipgram

Architecture: CBOW

CBOW has a similar architecture to skipgram however the input and output are flipped

Instead, context words are taken as an input, averaged and then the softmax is calculated for the target (center) word

Summary
- Input: Target word
- Output: Context words
- Output layer: softmax
- Loss: Cross entropy

Continuous Bag of Words (CBOW)

Several approaches to fitting Word2Vec models exist

	APPROACH	SPEED UP FACTOR	PERFORMANCE (SMALL VOCAB)	PERFORMANCE (LARGE VOCAB)
EXACT	Softmax	1x	Very good	Very poor
	Hierarchical Softmax	25x (50-100x)	Very poor	Very good
	Differentiated Softmax	2x	Very good	Very good
APPROXIMATE	CNN-Softmax	-	-	Bad-good
	Importance Sampling	(19x)	-	-
	Adaptive importance sampling	(100x)	-	-
	Target Sampling	2x	Good	Bad
	Noise Contrastive Estimation	8x (45x)	Very bad	Very bad
	Negative Sampling	(50-100x)	-	-
	Self-Normalisation	(15x)	-	-
	Infrequent Normalisation	6x (10x)	Very good	Good

Several variants exist to fit word embeddings.
Mikolov's seminal paper introduces negative sampling as an alternative to traditional softmax

Using embeddings

Why are neural embeddings preferred:
- Fast (compared to quadratic time scaling),
- Allow users to incorporate a new sentences/documents and
- Can easily add words to the vocabulary

They have two additional benefits:
1. Their performance for common linguistic tasks is superior to DSM methods and
2. They address non-linearities directly through their architecture rather than feature engineering

Practicalities:
- Embedding approaches do not only need to be with words
- [Facebook: Startspace - Embed All the Things](#)
- Evaluating embedding quality is challenging
- Embedding only makes sense when you have a lot of data and computational resources
- Most practitioners download pre-trained models and optimize for their network

Yoav Goldberg critique of word embeddings
- Mentions other classical methods that are tuned yield similar results
- Word embeddings are useful for general purpose situations however should not be viewed as the best approach for a particular language problem

Embeddings - what next?

A lot that wasn't covered

We didn't
1. Didn't deep dive into the math
2. Look at training the embedding
   - GloVe's SGD and ALS approaches are really interesting and build on earlier ideas
3. Discuss hyperparameter meanings or optimizations
4. Go into evaluation methods for embeddings
5. Briefly touched on different training methods but there's a lot more

References:

1. Evaluation Methods for Unsupervised Word Embedding - Tobias Schnabel, Igor Labutov, David Mimno, Thorsten Joachims (2015)
2. GloVe: Global Vectors for Word Representation - Jeffrey Pennington, Richard Socher, Christopher D. Manning (2014)
3. Efficient Estimation of Word Representations in Vector Space - Tomas Mikolov, Kai Chen, Greg Corrado, Jeffrey Dean (09/07/2013)

Course notes with more detail can be found at:

Word Embeddings 1- Intro and Background
(http://joshuahtouyz.com/blog/p1_word_embeddings_introduction)

Word Embeddings 2- Singular Value Decomposition
(http://joshuahtouyz.com/blog/p2_word_embeddings_svd)

Word Embeddings 3- Word2Vec
(http://joshuahtouyz.com/blog/p3_word_embeddings_word2vec_glove)

Word Embeddings 4- Evaluating Embeddings
(http://joshuahtouyz.com/blog/p4_embedding_evaluation_methods)